

DBMS

Introduction:

- Relation is described using 2 things:
 - i) Relation schema; ii) Relation instance
- Degree/arity; Cardinality:
- IC: PK; UNIQUE; NOT NULL; Check; Foreign key
- Key/Constraints: key/ck; sk; pk;
- FK constraints: referencing relation - child; referred relation - parents;
 - Every value under FK must be under PK.
 - violations: deletion/updation from parent insertion in child.
 - on del cascade; one update cascade;
 - set null on deletion; set default on del;

ER-Model:

- entity, entity set; relationship, relationship set;


Attributes

- simple, composite
- single val, Multi val
- stored, derived
- key
- Descriptive attribute

- Relationship Degree: no of entities in a relationship
 - Unary/recursive; binary; ternary; n-ary

ER-to-Relational:

- Composite att - Create multiple columns
- Multi val att - Create 2 tables
(Pk + non-multi-val) (Pk + Multi-val)
- For converting a relationship (□-◇-□) create 3 tables.
- key constraint: Entity that uniquely identifies relationship (→, ⇒)
- total participation: (⇒)
- atmost one; exactly one; atleast one; 0 or more
(→) (⇒) (⇒) (→)
- Cardinality ratios:




Weak Entity Set:

- weak entity set: no sufficient att to form key
- forms key when combined (partial key + Pk of strong entity)
- 2 tables
- FK on del cascade (think y)

Class hierarchy

- 3 tables without covering constraint
- 2 tables with covering constraint if overlapping is present anomalies



Aggregation: relationship over relationship

- 5 tables req;



Functional Dependencies:

- Trivial FD; Non-trivial FD;
- Properties:
 - Reflexivity: $X \rightarrow Y, Y \subseteq X$
 - Augmentation: $X \rightarrow Y \Rightarrow XZ \rightarrow YZ$
 - Transitivity: $X \rightarrow Y \wedge Y \rightarrow Z \Rightarrow X \rightarrow Z$
 - Union: $X \rightarrow Y \wedge X \rightarrow Z \Rightarrow X \rightarrow YZ$
 - Decomposition: $X \rightarrow YZ \Rightarrow X \rightarrow Y \wedge X \rightarrow Z$

- closure set of FD: F^+

- closure set of attributes

- Additional FDs can be determined from att. closure.

- Two FDs are said to be equivalent iff

$$(F^+ \subseteq G^+) \text{ and } (G^+ \subseteq F^+)$$

This is similar to $(F \text{ covers } G) \text{ and } (G \text{ covers } F)$

Minimal cover:

- Minimize LHS: In $AB \rightarrow C$, we can replace it with $A \rightarrow C$ if $BE \rightarrow C$ if $A \in B^+$

- Remove redundant productions. by checking if an FD can be obtained from remaining FDs.

Finding keys:

- * No of sks can be found using sum rule.
- * Independent att must be present in ck
- If Ax is ck and $y \rightarrow A$ is an FD then yx is also a ck.

Properties of Decomposition:

(i) Loss-less:

If original relationship can be obtained from child relations then it is lossless.

$R_1 \cap R_2$ must be key of R_1 or R_2



(ii) Dependency Preserving:

$$F^+ = (F_1 \cup F_2)^+$$

Normalization

~~only~~

→ If $X \rightarrow A$ is dependency and A is not prime attribute
if $x \in CK$ then $X \rightarrow A$ is Partial FD
else if A is not prime att and $x \notin CK$
then $X \rightarrow A$ is transitive FD

1NF: No multivalued att allowed. Only atomic values.

2NF: PFDs not allowed. Only full FDs allowed

3NF: Transitive FDs not allowed.

Every $X \rightarrow A$ (non-trivial) is
 X is SK (or) A is prime att.

BCNF: Prime transitivity not allowed.

For every $X \rightarrow A$, X is SK

3NF Decomposition:

- Find minimal cover of F
- Create a relation for each FD
- If no relation has key, then create a relation with key.
- Remove redundant relations.

BCNF Decomposition:

If $X \rightarrow A$ is dependency that violates BCNF property
create two relations $R_1(R-A)$ & $R_2(XA)$ and
continue this process until all the smaller relations
are in BCNF.

→ BCNF decomposition may not be dependency preserving.

→ Every all key relation is in BCNF.

Relational Algebra: Procedural; gives step by step process.

→ RA & RC eliminates duplicate tuples by default.

→ $\Pi_{att}(\sigma_{cond} table_name)$;

→ $\cup, \cap, -$ {relations must be compatible}

↳ same col & domain

→ Cross product;

→ Join

Conditional join: $R \bowtie S \cong \sigma_C(R \times S)$

Equi join: $R \bowtie S$ & C is equality condition (connected by \wedge)

Natural join: Equi join in which equality is implied.

→ Division: A/B is largest relation instance Q such that
 $Q \times B \subseteq A$ (A has 2 att, B has 1 att)

→ Rename: $\rho(R(F), E)$; $E \rightarrow$ relation expression.

$R \rightarrow$ new name; $F \rightarrow$ list of names & att (optional)

Relational Calculus: Non-procedural / Declarative. 115

TRC: $\{T/P(T)\}$; $T \rightarrow$ tuple variable.

DRC: $\{Q(d_1, d_2, \dots, d_n)/P(d_1, d_2, \dots, d_n)\}$; $d_i, d_n \rightarrow$ domain variables

→ \forall, \exists and ~~and~~ connectives ($\wedge, \vee, \Rightarrow, \Leftrightarrow$) can be used within expression P .

→ unsafe queries are possible within relational calculus.

→ Relational algebra \cong safe relational calculus

SQL:

→ Order: from, where, groupby, having, select expressions, distinct, order by.

→ string enclosed within single quotes. inclusive

→ operators: (and, or, not) between and & silly with not

is null is not null like: $=, <, >, <=, >=, <>$

in, not in exists(): true for non empty set

also works with multiple attributes (multiple supplies)

unique(): true when all values are distinct

any() & all() Eg: $X < \text{any}(2, 3, 4)$ & $X < \text{all}(2, 3, 4)$

any(\emptyset) = false; all(\emptyset) = true;

union; intersect; except; (sets must be compatible)

join: from table1 join table2 on condition

natural join: from table1 natural join table2

outer join: used to join all rows even if no matching row

left outer join (left join)

right outer join (right join)

full outer join (full join)

usage: T_1 left join T_2 on
(or)
 T_1 natural left join T_2

note: from T_1, T_2

where $T_1.att = T_2.att(+)$ → left join. Silly we have right join

Clauses:

WHERE: uses above operators for row selection.

* conditions in where clause should not have aggregate func.

Groupby & Having: group by att-list having condition;

→ select & having must contain either aggregate functions

or attributes that appeared in group by

→ without group by whole table is considered a single group

Orderby: order by col1 desc/asc col2 desc/asc...

→ Think how it works when sorted with multiple attributes

with: used for creating temporary relations.

Syntax: with table.name(att.list) as (expression)

table.name(att.list) as (expression) ...

Aggregate func: MIN(), MAX(), COUNT(), avg(), sum()

→ MIN, MAX, COUNT works with both numerical & strings

→ avg & sum only with numerical. → COUNT, SUM, AVG can appear with distinct.

→ at agg. func ignore null values.

→ $\text{avg}(c) = \text{sum}(c)/\text{count}(c)$; $\text{avg}(dis) = \text{sum}(dis)/\text{count}(dis)$

Locking Protocols

- shared lock, Exclusive lock

2PL: Growing phase & shrinking phase must be present in 2PL.

- Dead locks are possible
- If no deadlock then ensures CS.
- Edge of precedence graphs are drawn for conflicting locks
- Cycle in precedence graph \Leftrightarrow Deadlock.
- Cascading rollbacks & irrecoverable schedules are possible.

order of lock points determining order of serial

strict 2PL

- All exclusive locks are released only after commit operation.
- There will be no ww and wr conflicts

Both strict 2PL & rigorous 2PL are free from cascading rollbacks and hence recoverable. but deadlocks possible

Conservative 2PL:

- All the locks are held at once and released after commit
- This is free from cascading rollbacks, recoverable & deadlock free

Rigorous 2PL

- All locks are released only after commit operation
- No ww, wr, rw conflicts

Deadlock Handling:

(i) Deadlock prevention:

- wound-wait: older wounds, younger waits
- wait-die: older waits, young dies

- In both, unnecessary rollbacks are possible.
- In both, restarted T_i is started with same time stamp.

(ii) Deadlock Detection:

- Here a wait-for graph is maintained.
- Whenever a conflicting lock is req, an edge is added and removed when lock is acquired.
- Deadlock \Leftrightarrow cycle.

Multiple Granularity:

- If a child is locked, its ancestor can't be locked in conflicting mode. E.g. if parent is locked child can be locked in conflicting mode.
- locks: S, X, IS, IX, SIX
- Multiple granularity 2PL: locking: root to leaf, unlocking: leaf to root.

If $TS(T) < WTS(A)$ & $TS(T) \geq RTS(A)$ then it is not possible under TOP but possible under TWR

only if $TS(T) \geq WTS(A)$ and $TS(T) < RTS(A)$

Timestamp based protocols:

- Every schedule is under this protocol is equivalent to serial schedule T_i, T_j such that $TS(T_i) < TS(T_j)$
- This is deadlock free.

2 types:

(i) Time stamp Ordering Protocol:

- Here all conflicting operations are executed in the order of their time stamp.
- If an conflicting operation is performed out of order then that T_i will be aborted and rolled back and restarted with new time stamp.
- Starvation is possible.
- Ensures CS & deadlock free.

(ii) Thomas write rule:

- This is modification of TOP that minimizes the problem of starvation by avoiding obsolete write.
- i.e. writes conflicting with writes of new T_i .
- $w_i(A)$ is called obsolete write, if $w_i(A)$ occurs after $w_j(A)$ and $T_i < T_j$
- So here we just ignore $w_i(A)$
- TWR ensures VS.

- Cascading rollbacks & irrecoverable schedules are possible under TOP & TWR.



Implementation of TOP:

- Every T_i has timestamp $TS(T_i)$
- Every data item A is given two timestamps (read TS & write TS)
- $RTS(A)$: highest TS of any T_i that performed read on A.
- $WTS(A)$: highest TS of any T_i that performed write on A.

Read operation on A by T_i :

If $TS(T) \geq WTS(A)$ then read can be performed otherwise abort T_i and restart with new TS

write operation on A by T_i :

If $TS(T) \geq WTS(A)$ & $TS(T) \geq RTS(A)$ then write can be performed on A otherwise abort T and restart with new TS